



Markov Chains and Reward Process

Easwar Subramanian

TCS Innovation Labs, Hyderabad

Email: easwar.subramanian@tcs.com / cs5500.2020@iith.ac.in

August 5, 2022

Administrivia



- ▶ Please consult Prof. Vineeth, for all queries related to registration and other administrative issues.
- ▶ If need be, register for CS 5500 instead of AI 3000 (relevant for CS, PhD students).

Overview



- Review
- 2 Mathematical Framework for Decision Making
- Markov Chains
- Markov Reward Process



Review



Types of Learning: Summary



- Labeled data
- · Direct feedback
- · Predict outcome/future



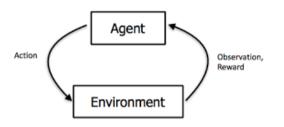
- · No labels
- · No feedback
- · "Find hidden structure"

- · Decision process
- · Reward system
- · Learn series of actions



Characteristics of Reinforcement Learning



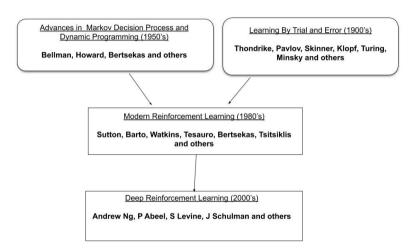


- ▶ Observations are <u>non i.i.d</u> and are sequential in nature
- ▶ Agent's action (may) affect the subsequent observation seen
- ▶ There is no supervisor; Only reward signal (feedback)
- ▶ Reward or feedback can be delayed



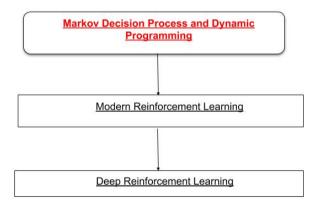
Reinforcement Learning: History





Course Setup





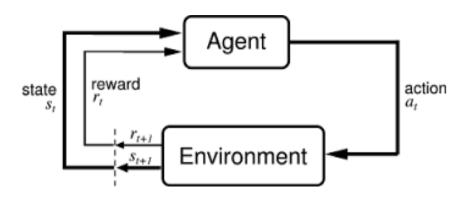




Mathematical Framework for Decision Making

RL Framework: Notations





Markov Decision Process



- ▶ Markov Decision Process (MDP) provides a <u>mathematical framework</u> for modeling decision making process
- ▶ Can formally describe the working of the environment and agent in the RL setting
- ➤ Can handle huge variety of interesting settings
 - ★ Multi-arm Bandits Single state MDPs
 - ★ Optimal Control Continuous MDPs
- ► Core problem in solving an MDP is to find an 'optimal' policy (or behaviour) for the decision maker (agent) in order to maximize the total future reward





Markov Chains



Random Variables and Stochastic Process



Random Variable (Non-mathematical definition)

A random variable is a variable whose value depend on the outcome of a random phenomenon

- Outcome of a coin toss
- ▶ Outcome of roll of a dice

Stochastic Process

A stochastic or random process, denoted by $\{s_t\}_{t\in T}$, can be defined as a collection of random variables that is indexed by some mathematical set T

- ▶ Index set has the interpretation of time
- ▶ The set T is, typically, \mathbb{N} or \mathbb{R}



Notations



- \blacktriangleright Typically, in optimal control problems, the index set is continuous (say \mathbb{R})
- ▶ Throughout this course (RL), the index set is always discrete (say \mathbb{N})
- ▶ Let $\{s_t\}_{t\in T}$ be a stochastic process
- ▶ Let s_t be the state at time t of the stochastic process $\{s_t\}_{t \in T}$

Markov Property



Markov Property

A state s_t of a stochastic process $\{s_t\}_{t\in T}$ is said to have Markov property if

$$P(s_{t+1}|s_t) = P(s_{t+1}|s_1, \cdots, s_t)$$

The state s_t at time t captures all relevant information from history and is a sufficient statistic of the future

Transition Probability



State Transition Probability

For a Markov state s and a successor state s', the state transition probability is defined by

$$\mathcal{P}_{ss'} = P(s_{t+1} = s' | s_t = s)$$

State transition matrix \mathcal{P} then denotes the transition probabilities from all states s to all successor states s' (with each row summing to 1)

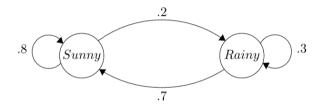
$$\mathcal{P} = \begin{bmatrix} \mathcal{P}_{11} & \mathcal{P}_{12} & \cdots & \mathcal{P}_{1n} \\ \vdots & & & \\ \mathcal{P}_{n1} & \mathcal{P}_{n2} & \cdots & \mathcal{P}_{nn} \end{bmatrix}$$

Markov Chain



A stochastic process $\{s_t\}_{t\in T}$ is a Markov process or Markov Chain if the sequence of random states satisfy the Markov property. It is represented by tuple $\langle S, P \rangle$ where S denote the set of states and P denote the state transition probability

Example 1: Simple Two State Markov Chain



- ightharpoonup State $S = \{Sunny, Rainy\}$
- ► Transition Probability Matrix

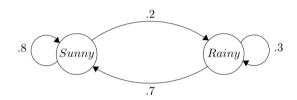
$$\mathcal{P} = \left[\begin{array}{cc} .8 & .2 \\ .7 & .3 \end{array} \right]$$





Markov Chain: Example Revisited





State $S = \{Sunny, Rainy\}$ and Transition Probability Matrix

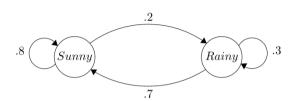
$$\mathcal{P} = \left[\begin{array}{cc} .8 & .2 \\ .7 & .3 \end{array} \right]$$

▶ Probability that tomorrow will be 'Rainy' given today is 'Sunny' = 0.2

Figure Source:

Multi-Step Transitions





Probability that day-after-tomorrow will be 'Rainy' given today is 'Sunny' is given by 0.2 * 0.3 + 0.8 * 0.2 = 0.22

In general, if one step transition matrix is given by,

$$\mathcal{P} = \left[\begin{array}{cc} P_{ss} & P_{sr} \\ P_{rs} & P_{rr} \end{array} \right]$$

then the two step transition matrix is given by,

$$\mathcal{P}_{(2)} = \left[\begin{array}{cc} P_{ss} * P_{ss} + P_{sr} * P_{rs} & P_{ss} * P_{sr} + P_{sr} * P_{rr} \\ P_{rr} * P_{rs} + P_{rs} * P_{ss} & P_{rr} * P_{rr} + P_{rs} * P_{sr} \end{array} \right] = P^2$$



Figure Source:

Multi-Step Transitions



In general, n-step transition matrix is given by,

$$P_{(n)} = P^n$$

Assumption

We made an important assumption in arriving at the above expression. That the one-step transition matrix stays constant through time or independent of time

- Markov chains generated using such transition matrices are called <u>homogeneous</u>
 Markov chains
- ▶ For much of this course, we will consider homogeneous Markov chains, for which the transition probabilities depend on the length of time interval $[t_1, t_2]$ but not on the exact time instants

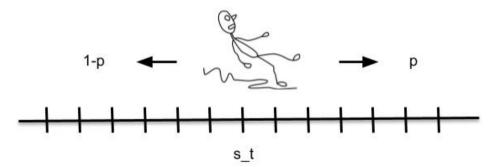
Markov Chains: Examples



Example 2: One dimensional random walk

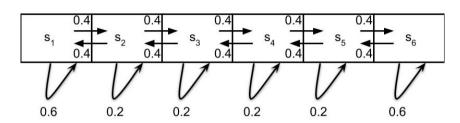
A walker flips a coin every time slot to decide which 'way' to go.

$$s_{t+1} = \begin{cases} s_t + 1 & \text{with probability } p \\ s_t - 1 & \text{with probability } 1 - p \end{cases}$$



Example 3: Simple Grid World





- $\triangleright \mathcal{S} = \{s_1, s_2, s_3, s_4, s_5, s_6, s_6\}$
- $\triangleright \mathcal{P}$ as shown above
- ightharpoonup Example Markov Chains with s_2 as start state

$$\star \{s_2, s_3, s_2, s_1, s_2, \cdots\}$$

$$\star \{s_2, s_2, s_3, s_4, s_3, \cdots\}$$



Markov Chains: Examples



Example 4: Dice roll experiment

Let $\{s_t\}_{t\in T}$ model the stochastic process representing the cumulative sum of a fair six-sided die rolls

Example 5: Natural Language Processing

Let $\{s_t\}_{t\in T}$ model the stochastic process that keeps track of the chain of letters in a sentence. Consider an example

Tomorrow is a sunny day

- ▶ We normally don't ask the question what is probability of character 'a' appearing given previous character is 'd'
- ► Sentence formation is typically **non-Markovian**



Notion of Absorbing State



Absorbing State

A state $s \in \mathcal{S}$ is called **absorbing** state if it is impossible to leave the state. That is,

$$P_{ss'} = \left\{ \begin{array}{ll} 1, & \text{if } s = s' \\ 0, & \text{otherwise} \end{array} \right\}$$



Markov Reward Process



Markov Reward Process



Markov Reward Process

A Markov reward process is a tuple $\langle S, \mathcal{P}, \mathcal{R}, \gamma \rangle$ is a Markov chain with values

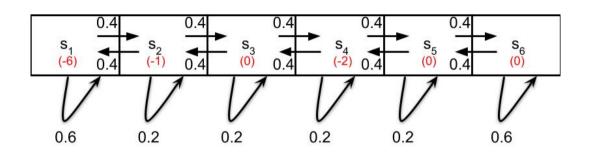
- \triangleright S: (Finite) set of states
- $\triangleright \mathcal{P}$: State transition probablity
- \triangleright \mathcal{R} : Reward for being in state s_t is given by a deterministic function \mathcal{R}

$$r_{t+1} = \mathcal{R}(s_t)$$

 $ightharpoonup \gamma$: Discount factor such that $\gamma \in [0,1]$

Simple Grid World: Revisited





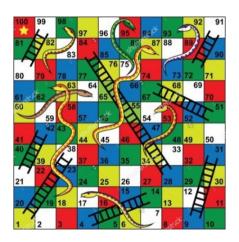
▶ For the Markov chain $\{s_2, s_3, s_2, s_1, s_2, \cdots\}$ the corresponding reward sequence is $\{-1,0,-1,-6,-1,\cdots\}$

No notion of action



Example: Snakes and Ladders





Example: Snakes and Ladders



- ▶ States $S: \{s_1, s_2, \cdots, s_{100}\}$
- ▶ Transition Probability P:
 - ★ What is the probability to move from state 2 to 6 in one step?
 - ★ What are the states that can be visited in one-step from state 2?
 - \star What is the probability to move from state 2 to 4?
 - \star Can we transition from state 15 to 7 in one step?

Question: Is transition matrix independent of time?

Question: Can we formulate the game of Snake and Ladders as a MRP?

Need to define suitable reward function and discounting factor



On Rewards: Total Return



- At each time step t, there is a reward r_{t+1} associated with being in state s_t
- ▶ Ideally, we would like the agent to pick such trajectories in which the cumulative reward accumulated by traversing such a path is high

Question: How can we formalize this?

Answer: If the reward sequence is given by $\{r_{t+1}, r_{t+2}, r_{t+3}, \cdots\}$, then, we want to maximize the sum

$$r_{t+1} + r_{t+2} + r_{t+3} + \cdots$$

Define G_t to be

$$G_t = r_{t+1} + r_{t+2} + r_{t+3} + \dots = \sum_{k=0}^{\infty} r_{t+k+1}$$

The goal of the agent is to pick such paths that maximize G_t



Total (Discounted) Return



Recall that,

$$G_t = r_{t+1} + r_{t+2} + r_{t+3} + \dots = \sum_{k=0}^{\infty} r_{t+k+1}$$

▶ In the case that the underlying stochastic process has infinite terms the above summation could be divergent

Therefore, we introduce discount factor $\gamma \in [0,1]$ and redefine G_t as

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

- \triangleright G_t is the total discounted return starting from time t
- ▶ If γ < 1 then the infinite sum has a finite value if the reward sequence is bounded
- \blacktriangleright γ close to 0 the agent is concerned only with immediate reward(s) (myopic)
- ightharpoonup close to 1 the agent considers future reward more strongly (far-sighted)



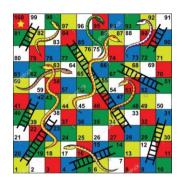
Few Remarks on Discounting



- ▶ Mathematically convienient to discount rewards
- ▶ Avoids infinite returns in cyclic and infinite horizon setting
- ▶ Discount rate determines the present value of future reward
- ▶ Offers trade-off between being 'myopic' and 'far-sighted' reward
- ▶ In finite MDPs, it is sometimes possible to use undiscounted reward (i.e. $\gamma = 1$), for example, if all sequences terminate

Snakes and Ladders: Revisited





Question: What can be a suitable reward function and discount factor to describe 'Snake and Ladders' as a Markov reward process?

- ▶ Goal: From any given state reach s_{100} in as few steps as possible
- ▶ Reward $\mathcal{R}: \mathcal{R}(s) = -1$ for $s \in s_1, \dots, s_{99}$ and for $R(s_{100}) = 0$
- ▶ Discount Factor $\gamma = 1$

